

Bayesian parameter synthesis of Markov population models

Master Project on Modeling of Complex, Self-organizing
systems

Huy Phung
University of Konstanz

May 1, 2020

Abstract

We study the collective behavior of a bee colony. Each bee in a colony could sting after observing a threat in the surrounding environment, and warn other bees by releasing a special substance, pheromone. By sensing the pheromone released in the environment, other bees in the colony may also sting. However, since stinging leads to the termination of an individual bee, it reduces the total defense capability as well. The stinging behavior of bee colonies can be studied by observing the changes in the bee population. In this project, we use parametric Markov Chain to model the bee population. The contributions of this project are (1) to show that steady-state distribution can be more precisely estimated with fewer data using Bayesian inference, and (2) Bayesian inference can also be used to estimate model parameters of parametric Markov Chain with more scalability (lower computational cost for models with a higher number of parameters, compare with the existing method with frequentist approach).

1 Preliminaries

1.1 Discrete time Markov Chain

Assume that each bee in a colony decides its next action (to sting or not to sting) based only on the current state of the environment, and the number of bees who sting or not sting can be modeled as a Markov process. To reduce the complexity of the model, we make another assumption that the observations on a bee colony are conducted in uniform time duration, hence the model is of discrete-time. Therefore, Discrete-Time Markov Chain is the selected mathematical model to describe the bee colony population. The definitions and notations on this section are based on [1] and [3]

Definition 1.1 (Discrete Time Markov Chain). A Discrete Time Markov Chain (DTMC) is a tuple $(S, \mathbf{P}, S_{init}, AP, L)$

- S is a countable nonempty set of *states*
- $\mathbf{P} : S \times S \rightarrow [0, 1]$ is the *transition probability* function, s.t

$$\sum_{s' \in S} \mathbf{P}(s, s') = 1$$

- $S_{init} : S \rightarrow [0, 1]$ is the *initial distribution*, s.t

$$\sum_{s' \in S} S_{init}(s') = 1$$

- AP is a set of *atomic propositions*
- $L : S \rightarrow 2^{AP}$ is the labelling function on states.

The transition probability function in DTMC is given as a stochastic matrix \mathbf{P} , that is, \mathbf{P} satisfies the following properties

- \mathbf{P} is a square matrix.
- all elements are in $[0, 1]$.
- row elements sum up to 1.

Given a model \mathcal{M} , the population in steady state is represented by a set of special states, namely *terminal Strongly Connected Components*

Definition 1.2 (Terminal Strongly Connected Components). Let $\mathcal{M} = (S, \mathbf{P}, S_{init}, AP, L)$ a DTMC, a state $s \in S$ is a Terminal Strongly Connected Component (tSCC for short) if and only if $P(s, s) = 1$ and $P(s, s') = 0 \forall s' \in S, s' \neq s$

In order to generalize the stochastic matrix to encompass unknown information of the system, we introduce *parametric DTMC*. Let $\theta = (\theta_1, \dots, \theta_n) \in [0, 1]^n$, we represent each element in matrix \mathbf{P} as a polynomial function of θ . Let \mathbf{Pol}_θ be the set of polynomials $P : [0, 1]^n \rightarrow [0, 1]$. We define *parametric Discrete-Time Markov Chain*, or *pMC* for short, as follow:

Definition 1.3 (Parametric DTMC). A parametric Discrete Time Markov Chain (pMC for short) is a tuple $(S, \theta, \mathbf{P}_\theta, S_{init}, AP, L)$

- S is a countable nonempty set of *states*
- θ is the set of model parameters
- $\mathbf{P} : S \times S \rightarrow \mathbf{Pol}_\theta$ is the *transition probability* function that map a transition relation between two states to a polynomial function of θ
- $S_{init} : S \rightarrow [0, 1]$ is the *initial distribution*, s.t

$$\sum_{s' \in S} S_{init}(s') = 1$$

- AP is a set of *atomic propositions*
- $L : S \rightarrow 2^{AP}$ is the labelling function on states.

A concrete assignment of θ on pMC induces a DTMC. Let ϕ be a specification on a pMC \mathcal{M} (for example, in this project ϕ is the steady-state distribution), we estimate the model parameter θ , such that $\mathcal{M} \models \phi$

1.2 Bayesian inference

1.2.1 Bayesian parameter estimation

Let D be observed data. In statistical inference, we assume that the observed data has a probability distribution of unknown parameter θ , i.e $D \sim P(D|\theta)$.

In frequentist approach, the estimation of θ based on long-run property, that is, given a large enough sample size, expected value of parameter estimation $\hat{\theta}$ is equal to θ . Therefore, frequentist approach requires to gather a large amount of data to deliver a close estimation $\hat{\theta}$. In Bayesian approach, we reuse the information *beliefs* gained from observed data to enhance the accuracy of the estimation of $\hat{\theta}$. The main advantage of Bayesian approach over frequentist approach is that it require less data to obtain an estimation $\hat{\theta}$. The beliefs obtained from prior knowledge of model parameter θ is represented by *prior distribution* $\pi(\theta)$.

Also, we have probability distribution of observed data, given parameter θ , $P(D|\theta)$. This is also called *likelihood function*.

With Bayesian formula, we have

$$\pi(\theta|D) = \frac{P(D|\theta)\pi(\theta)}{\int_{\theta} P(D|\theta)\pi(\theta)d\theta}$$

$\int_{\theta} P(D|\theta)\pi(\theta)d\theta$ is called *marginal distribution*. $\pi(\theta|D)$ is called *posterior distribution*. Computing posterior distribution is the essential part of Bayesian inference, since it gives us the estimation of parameter θ .

1.2.2 Posterior conjugation

Conjugated posteriors are special cases of Bayesian inference, in which the prior and posterior distribution belongs to the same family of distribution. Conjugated posteriors give us significant benefits

1. Tractability: we have analytical form of posterior distribution.
2. Computationally effective: updating model parameter is of linear time to the dimension of parameter.

We consider two conjugated posterior: Binomial-Beta and Dirichlet-Multinomial

Lemma 1 (Binomial-Beta Conjugation). Binomial distribution is conjugated to beta distribution.

Proof. The observed data $D = (x_1, \dots, x_n)$ is sampled from *Binomial*(k, θ) function

$$P(D|\theta) = \prod_{i=1}^n \binom{k}{x_i} \theta^{x_i} (1 - \theta)^{k-x_i}$$

The parameter θ is of $Beta(\alpha, \beta)$ distribution

$$\pi(\theta) = \theta^{\alpha-1}(1-\theta)^{\beta-1}$$

We obtained:

$$\begin{aligned}\pi(\theta|D) &\sim P(D|\theta)\pi(\theta) \\ &\sim \theta^{\sum_{i=1}^n x_i} (1-\theta)^{nk - \sum_{i=1}^n x_i} \theta^{\alpha-1} (1-\theta)^{\beta-1} \\ &= \theta^{\alpha-1 + \sum_{i=1}^n x_i} (1-\theta)^{\beta-1 + nk - \sum_{i=1}^n x_i}\end{aligned}$$

Thus, the posterior is $Beta(\alpha + \sum_{i=1}^n x_i, \beta + nk - \sum_{i=1}^n x_i)$ \square

Generalize this conjugation, we also have Multinomial-Dirichlet conjugation.

Lemma 2 (Multinomial-Dirichlet Conjugation). Multinomial distribution is conjugated to Dirichlet distribution.

Proof. The observed data $D = (x_1, \dots, x_n)$ is sampled from $Multinomial(n; \theta_1, \dots, \theta_n)$ function

$$P(x_1, \dots, x_n | N, \theta_1, \dots, \theta_n) = \frac{n!}{x_1! \dots x_n!} \prod_{i=1}^n \theta_i^{x_i}$$

The parameter $(\theta_1, \dots, \theta_n)$ is $Dirichlet(\alpha_1, \dots, \alpha_n)$

$$\pi(\theta_1, \dots, \theta_n) = \frac{1}{\mathbf{B}(\alpha_1, \dots, \alpha_n)} \prod_{i=1}^n \theta_i^{\alpha_i-1}$$

We obtain

$$\begin{aligned}\pi(\theta_1, \dots, \theta_n | D) &\sim P(D|\theta)\pi(\theta) \\ &\sim \prod_{i=1}^n \theta_i^{x_i} \prod_{i=1}^n \theta_i^{\alpha_i-1} \\ &\sim \prod_{i=1}^n \theta_i^{\alpha_i-1 + \sum_{i=1}^n x_i}\end{aligned}$$

Thus, the posterior is $Dirichlet(\alpha_1 + x_1, \dots, \alpha_n + x_n)$ \square

More detailed description in these cases can be found in [6] and [2]. We summarize the necessary results in the following table:

Likelihood	Prior	Posterior parameters
$Binomial(k, \theta)$	$Beta(\alpha, \beta)$	$\alpha' = \alpha + \sum_{i=1}^n x_i, \beta' = \beta + nk - \sum_{i=1}^n x_i$
$Multinomial(n; \theta_1, \dots, \theta_n)$	$Dirichlet(\alpha_1, \dots, \alpha_n)$	$\alpha'_i = \alpha_i + x_i, 1 \leq i \leq n$

1.2.3 Markov Chain Monte Carlo method

In case the posterior distribution has no analytical form or its analytical form is difficult to sample from directly, we need to use *Metropolis-Hastings* algorithm (MH in short). The reason of using MH algorithm to MH algorithm can draw samples from posterior distribution with only the likelihood function. Using the MH algorithm, we can estimate the parameter by posterior mean, without knowing the analytical form of posterior distribution itself.

Algorithm 1 Metropolis-Hastings Algorithm, D is the observation data, θ is parameter

```

1: procedure METROPOLIS-HASTINGS( $D$ , maxIteration)
2:   Select a proposal distribution  $\pi(\theta)$ 
3:   Initialize  $\theta$ 
4:   while maxIteration not reached do
5:      $L \leftarrow P(D|\theta)$ 
6:     Draw  $\theta'$  from the proposal distribution (transition kernel).
7:      $L' \leftarrow P(D|\theta')$ 
8:     if  $\ln(L') - \ln(L) > 0$  then
9:        $\theta = \theta'$  (accept  $\theta'$ )
10:    else
11:      with probability  $\epsilon$  accept  $\theta'$  (avoiding local maxima)
12:    end if
13:  end while
14: end procedure

```

The likelihood function can be implemented as log-likelihood to avoid underflow error. Proposal distribution defines how do we proceed to the next parameter value on the parameter space; it can be of any distribution family. In this project, since parameters are in $[0, 1]$, we select Beta distribution as proposal.

There are different methods to select the initial value of parameter P . In

this project, we initialize P randomly. There are two advantages of using Markov Chain Monte Carlo in Bayesian inference:

1. Parameter transition only needs the computation of likelihood function. Therefore, Monte Carlo Markov Chain can be used in general Bayesian inference, in which we are not guaranteed to have an analytical form of posterior.
2. Specifically in Metropolis-Hastings algorithm, marginal distribution is cancelled out, thus make MH a computationally efficient algorithm.

However, MH algorithm also has a drawback; its convergence becomes slower as the dimension of parameter θ increases.

1.2.4 Bayesian Parameter Estimation

With posterior distribution $\pi(\theta|D)$ we estimate the parameter $\hat{\theta}$ using Bayesian posterior mean:

$$\hat{\theta} = \mathbf{E}[\theta] = \int_{\theta} \theta \pi(\theta|D) d\theta$$

In case we have samples from posterior distribution, for example when we use MH algorithm, the discrete form of posterior mean is used:

$$\hat{\theta} = \mathbf{E}[\theta] = \sum_{\theta} \theta \pi(\theta|D)$$

Definition 1.4 (Bayesian Credible Set). Set C is a $(1-\alpha)100\%$ credible set for the parameter θ if the posterior probability for θ to belong to C equals $(1-\alpha)$.

$$P(\theta \in C|D) = \int_C \pi(\theta|D) d\theta = 1 - \alpha$$

In this project, we use 0.95 credible set, i.e $\alpha = 0.05$

Definition 1.5 (Highest Posterior Density credible set). Highest Posterior Density $(1 - \alpha)100\%$ credible set (HDP for short) is the interval with minimum length over all Bayesian $(1 - \alpha)100\%$ Credible Set.

In this project, the HDP is calculated using PyMC3 library [5]

2 Modeling of bees colony

2.1 Biological description

We consider the collective action of a bee colony. Each bee in a colony could possibly sting after observing a threat in the surrounding environment, and warn other bees by releasing pheromone. By sensing the pheromone released in the environment, other bees in the colony may also sting. Since stinging leads to the termination of an individual bee, it reduces the total defense capability as well. We studies how the actions of a bee changes with regarding to its surrounding the environment. There are 3 assumptions on the system:

1. Each bee release an unit amount of pheromone immediately after stinging.
2. A bee dies after stinging and releasing pheromone. In the other words, no bee can sting more than once.
3. Stinging behaviour only depends on the concentration of pheromone in the environment.

Under these assumption, a bee colony can be viewed as a set of agents (bees) interact with each other in a closed environment with the appearance of a factor *pheromone*. Each agent in the colony observe the following two factors of the environment: (1) amount of pheromone and (2) number of other agents. Afterward, the agent has probability to commit an action, namely *sting*. The agent is eliminated from environment after stinging.

2.2 Markov population models for bee colony

Semantics of Markov population models for bees colony are developed by [3]. In this report we summarize three models of bees population

1. Synchronous model
2. Asynchronous model
3. Semi-synchronous model

In the scope of this project we do not analyze the semantics of each model. Instead, the inference methods presented only uses the polynomial function of the models' steady-state distribution.

Asynchronous model

In asynchronous experiment we assume that there is almost improbable for two bees to sting simultaneously, and any bees release pheromone immediately after its death. By that observation we can assume that each bee sting at different level of pheromone.

Synchronous model

In *fully synchronous* experiment we assume that the number of stinging bees is only counted after a fixed amount of time, so that without loss of generality we can assume the pheromone diffuse almost immediately among the bee colony, and each bee decide to sting or not to sting immediately after sensing the pheromone concentration.

Semi-synchronous model

In *Semisynchronous model*, we assume that the behaviour is initially synchronous. That means, at the initial states we do synchronous update. From all succeeding states from initial states, the updates are of asynchronous semantics.

For a population of n bees, all three types of models share two properties:

1. Has exactly one initial state, $|S_{init}| = 1$. This assumption is obvious, since at the beginning of the experiment, all bees are alive.
2. Has $n + 1$ tSCCs ($tSCC_0, \dots, tSCC_n$). This is because the number of dead bees cannot exceed the population size. It also follows that

$$\sum_{i=0}^n P(FG \ tSCC_i) = 1$$

3 Bayesian parameter synthesis of Markov population model

3.1 Data

Let n be the population of bee in the colony. Assume that these n bees are in a closed environment, like a box. We use an initial stimulation, for example threatening the bees. Assume that all experiment have the same initial stimulation, we observe how many bees are dead in the steady state. This experiment has $n + 1$ outcomes $0, 1, \dots, n$ dead bees.

Assume we conduct N experiments, let x_0, x_1, \dots, x_N be the number of dead bees. Since in each experiment, we have $(0 \leq x_1, \dots, x_N \leq n + 1)$, we can place x_i on $n + 1$ bins from 0 to n . Assume we model the system by pMC \mathcal{M} with steady-state distribution $\theta_0, \dots, \theta_n$. Let a_i be the number of experiments which has outcome $i, 0 \leq i \leq n$, we obtain that a_i has *multinomial distribution*:

$$\begin{aligned} P(a_0, \dots, a_n | N, \theta_0, \dots, \theta_n) &\sim \text{Multinomial}(a_0, \dots, a_n | N, \theta_0, \dots, \theta_n) \\ &= \frac{n!}{a_0! \dots a_n!} \theta_0^{a_0} \dots \theta_n^{a_n} \end{aligned}$$

We use the multinomial distribution as likelihood function (data model) in the following Bayesian inferences.

3.2 Inference of steady state distribution

Using the Multinomial-Dirichlet conjugation, we can infer the steady state distribution as we observe new experiment data $D = (a_0, \dots, a_n)$, with a_i is the number of experiment in which the population in steady-state is i

Algorithm 2 Estimation of steady-state distribution given a sample S

- 1: **procedure** ESTIMATE $p_i = P(tSCC_i)(S)$
 - 2: Initialize $\alpha_0 = \alpha_1 = \dots = \alpha_n = 1$
 - 3: Initialize $p_0 = p_1 = \dots = p_n = 0$
 - 4: Update $\alpha_i = \alpha_i + a_i, 1 \leq i \leq n$
 - 5: Update $p_i = \alpha_i / \sum_{i=1}^n \alpha_i, 1 \leq i \leq n$
 - 6: **end procedure**
-

As later shown in the implementation with synthetic data, the Bayesian estimation of steady-state distribution converges quickly to the true parameter used for data synthesis. This method does not directly estimate the pMC model parameters. However, it can support the parameter estimation presented [3] to estimate the model parameters with narrower credible set.

3.3 Inference of model parameters

We can also use Bayesian inference to estimate pMC model parameters directly. Given a pMC \mathcal{M} with $\theta = (\theta_1, \dots, \theta_k)$ as its model parameters. Since $\theta_i \in [0, 1]$, we can assume that θ_i are of $Beta(\alpha, \beta)$ distribution. Note that α, β are hyperparameters and must be selected manually. Also note that $\theta_1, \dots, \theta_k$ are not simplex, thus it is not possible to use Dirichlet prior.

In this method, since there is no possible use of posterior conjugation, thus the hyperparameters estimation is hard and not in the scope of this project. Let $f_i = P(FG \text{ } tSCC_i)$ are polynomial functions of steady-state distribution, we denote $\epsilon_i = f_i(\theta)$ as evaluations of the polynomial function with a concrete assignment of θ . The steady-state distribution has then $Multinomial(N, \epsilon = (\epsilon_0, \dots, \epsilon_n))$ with N is the sample size of experiment data D . The posterior distribution has the following form:

$$\pi(\theta|D) \sim Multinomial(N, (\epsilon_0, \dots, \epsilon_n))\pi(\theta_1) \dots \pi(\theta_k)$$

As the posterior $\pi(\theta|D)$ has no analytical form, we use Metropolis-Hastings to sample from it.

Algorithm 3 Estimation of model parameters given a sample D

```
1: procedure ESTIMATE  $\theta(S)$ 
2:   Select hyperparameter  $\alpha, \beta$ 
3:   Select  $Beta(\alpha, \beta)$  as proposal distribution (transition kernel) for
    $\theta_i, 1 \leq i \leq n$ 
4:   Initialize  $\theta$ 
5:   while maxIteration not reached do
6:     Evaluate  $\epsilon = (f_0(\theta), \dots, f_n(\theta))$ 
7:      $L \leftarrow Multinomial(D|\epsilon)$ 
8:     Draw  $\theta'$  from the proposal distribution.
9:     Evaluate  $\epsilon' = (f_0(\theta'), \dots, f_n(\theta'))$ 
10:     $L' \leftarrow P(D|\epsilon')$ 
11:    if  $\ln(L') - \ln(L) > 0$  then
12:       $\theta = \theta'$  (accept  $\theta'$ )
13:    else
14:      with probability  $\epsilon$  accept  $\theta'$  (avoiding local maxima)
15:    end if
16:  end while
17: end procedure
```

3.4 Accuracy measures

In order to evaluate the accuracy of parameter point inference, we have to select a distance measure from estimated parameter to true parameter. In this project we use two measures:

Definition 3.1 (RMSE). Let $\theta = (\theta_1, \dots, \theta_n)$ and $\hat{\theta} = (\hat{\theta}_1, \dots, \hat{\theta}_n)$ be vectors of n real numbers, the Root Mean Square Error between p and \hat{p} is defined as follow:

$$RMSE(\theta, \hat{\theta}) = \sqrt{\frac{\sum_{i=1}^n (\theta_i - \hat{\theta}_i)^2}{n}}$$

4 Implementation and results

The Bayesian inference methods are implemented in Python 3 and tested in the followin system configuration:

- Intel Core i5-8265U @ 1.60GHz
- 16GB RAM
- OpenSUSE Tumbleweed 20200427
- Anaconda 3 2019.10 for linux x86_64

The following experiment on Bayesian inference use synthetic data from semi-synchronous models of 2, 5 and 10 bees (in order to compare the performance with related work in [3]).

4.1 Data synthesis

The data synthesis for n bees population is conducted as follow:

- Select a pMC model \mathcal{M} to model the population of . Let $\bar{p} = p_1, \dots, p_n$ is the model parameter vector.
- Assign a concrete value for \bar{p} , called \bar{p}_{true}
- From \mathcal{M} , deliver the rational functions of $f_i(\bar{p}) = P(FG \ tSCC_i)$, $0 \leq i \leq n$. (using PRISM Model Checker [4])
- Evaluate concrete value $\theta_i = f_i(\bar{p}_{true})$
- Draw sample S from multinomial distribution $Multinomial(N, \theta = (\theta_0, \dots, \theta_n))$
- S is the synthetic result for the experiment 3.1




Figure 1: Steady state data synthesis, semi-synchronous model of 5 bees population, model parameters (0.3, 0.25, 0.35, 0.45, 0.5)

4.2 Inference of steady state distribution

This experiment is conducted on semisynchronous model of 5 bees. The polynomial functions of steady-state distribution is generated by PRISM and parsed into Python 3 source code (`semisync_5bees.py`). The source code for this experiment is on Jupyter Notebook file `visualization.ipynb`




Figure 2: Bayesian estimation of steady state distribution

In order to compare with the work by [3], we also visualize the interval calculated by the following formula:

$$\theta_i$$




Figure 3: Comparison of Bayesian Highest Posterior Density 95% and interval estimated using method in [3]

4.3 Inference of model parameters

This experiment is conducted in 2, 5, and 10 bees population models of 2, 5, and 10 parameters respectively. All models are semi-synchronous.

```
Experiment with 2 bees , 2 params , semisync
Finished in 2.0310990789876087 seconds , chain length 50000
True parameter: [0.1 , 0.2]
Estimated parameter: [0.10183236 0.18434951]
Log likelihood: -8.4331030579051
AIC: 20.8662061158102
```

```
Experiment with 5 bees , 5 params , semisync
Finished in 7.2764098109910265 seconds , chain length 50000
True parameter: [0.1 , 0.2 , 0.4 , 0.5 , 0.6]
Estimated parameter: [
0.10135125 0.21855681 0.52070207 0.45761621 0.57840823
]
Log likelihood: -70.33078101358842
AIC: 150.66156202717684
```

```
Experiment with 10 bees , 10 params , semisync
Finished in 602.7453691669944 seconds , chain length 50000
```

```

True parameter: [0.1, 0.2, 0.4, 0.5, 0.6, 0.1, 0.2, 0.4, 0.5, 0.6]
Estimated parameter: [
0.11375671 0.14925181 0.3042601 0.57028791 0.81435131
0.10468808 0.37077345 0.40401714 0.28729648 0.28731649
]
Log likelihood: -395.266357301065
AIC: 810.53271460213

```

5 Conclusion

The goal of this project is to experiment Bayesian approach on parameter inference of Markov population models. The results shows that Bayesian inference is efficient to deliver an estimation of model parameters. Furthermore, the implementation is more efficient than the work in [3] as it can work with higher number of parameters and requires less computational resources. However, there are still problems

1. MH algorithm converges slower as the number of parameter increases, thus requires longer chain of sampling the parameter space, which may increase the computational cost dramatically.
2. PRISM is still necessary to deliver polynomial functions, thus introduce an overhead, that PRISM itself is resource-hungry.
3. In this project, there is no Bayesian inference, which is designed to deal with the case when the polynomial functions are hard to synthesize, or impossible to synthesize.

However, we believe that the method can be developed further to overcome these limitations.

asd

References

- [1] Christel Baier and Joost-Pieter Katoen. *Principles of model checking*. MIT press, 2008.
- [2] Michael Baron. *Probability and statistics for computer scientists*. CRC Press, 2019.

- [3] Matej Hajnal et al. “Data-Informed Parameter Synthesis for Population Markov Chains”. In: *International Workshop on Hybrid Systems Biology*. Springer. 2019, pp. 147–164.
- [4] M. Kwiatkowska, G. Norman, and D. Parker. “PRISM 4.0: Verification of Probabilistic Real-time Systems”. In: *Proc. 23rd International Conference on Computer Aided Verification (CAV’11)*. Ed. by G. Gopalakrishnan and S. Qadeer. Vol. 6806. LNCS. Springer, 2011, pp. 585–591.
- [5] John Salvatier, Thomas V Wiecki, and Christopher Fonnesbeck. “PyMC3: Python probabilistic programming framework”. In: *ascl* (2016), ascl–1610.
- [6] Stephen Tu. “The dirichlet-multinomial and dirichlet-categorical models for bayesian inference”. In: *Computer Science Division, UC Berkeley* (2014).